




Containers

Docker



Hypervisor and Virtual Machines

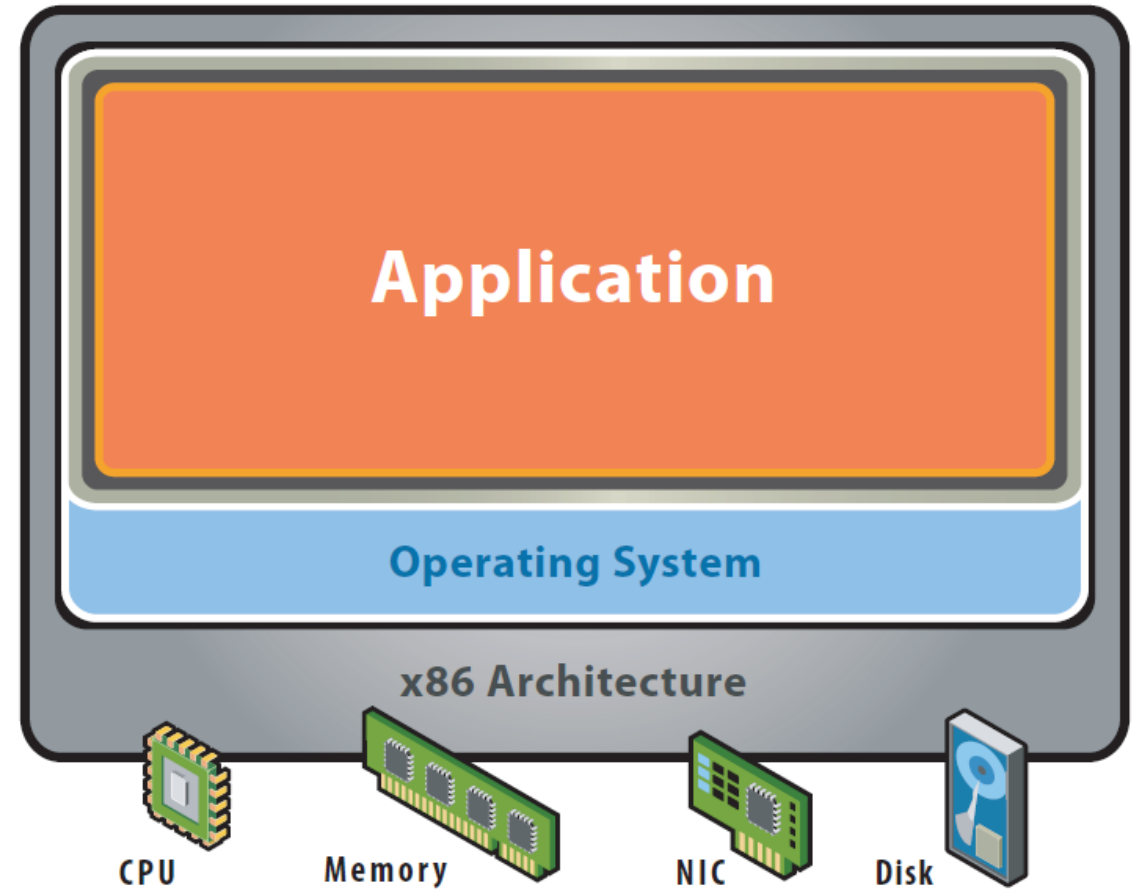
- Hypervisor is a process that separates the operating system from the hardware
 - Also known as Virtual Machine monitor (VMM)
- Virtual machines is a software that emulates a computer system





What is Virtualization?

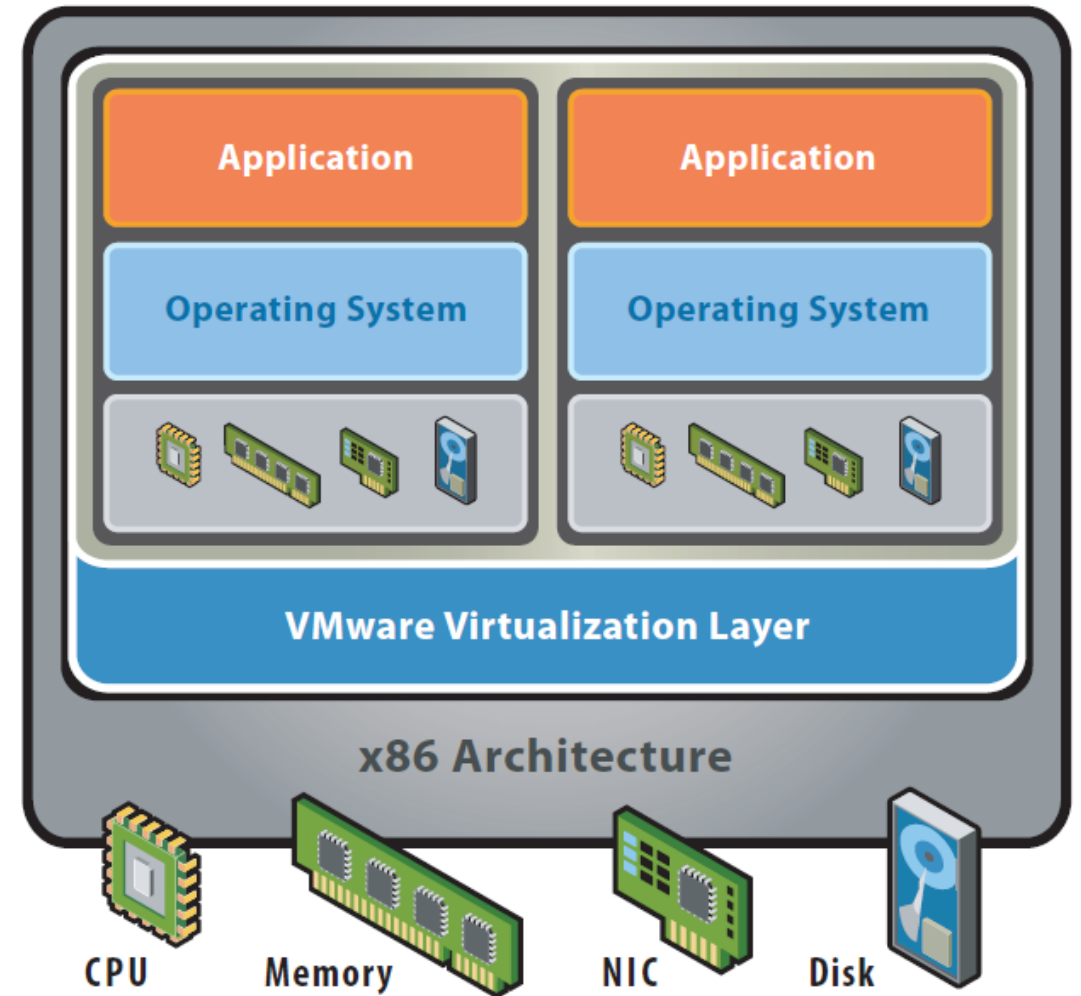
- Technology that allows you to run multiple operating systems on a single computer concurrently
 - Using hypervisor and virtual machines
- Typical computer setup
 - Single OS per machine
 - Run multiple applications on the same machine – may have dependencies conflict
 - Under utilize resources






Benefits of Virtualization

- Hardware independent from the operating system
- Manage OS and application as a single unit by encapsulating them in a virtual machine
- Increase the system's utilization
 - By allowing 2 or more operating system to run concurrently
- Distinct from dual boot
 - Only one operating system is active at anytime

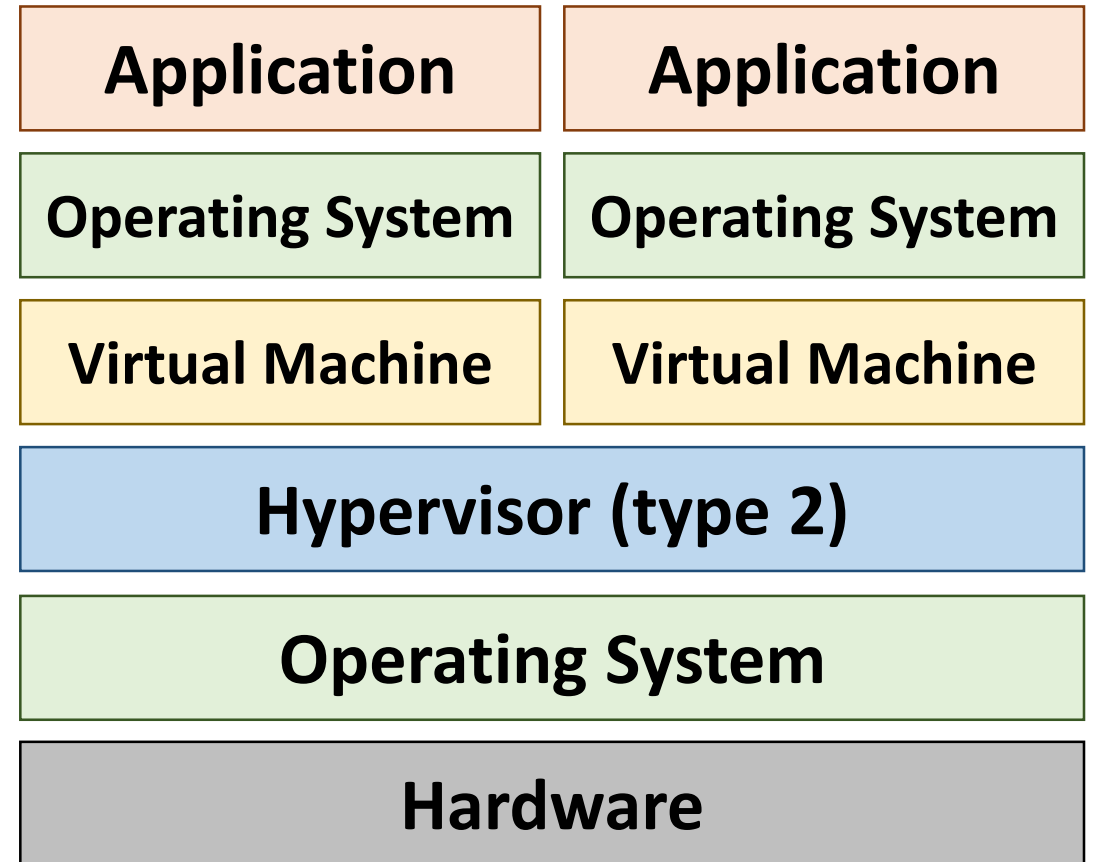




Types of Virtualization



Type 1



Type 2



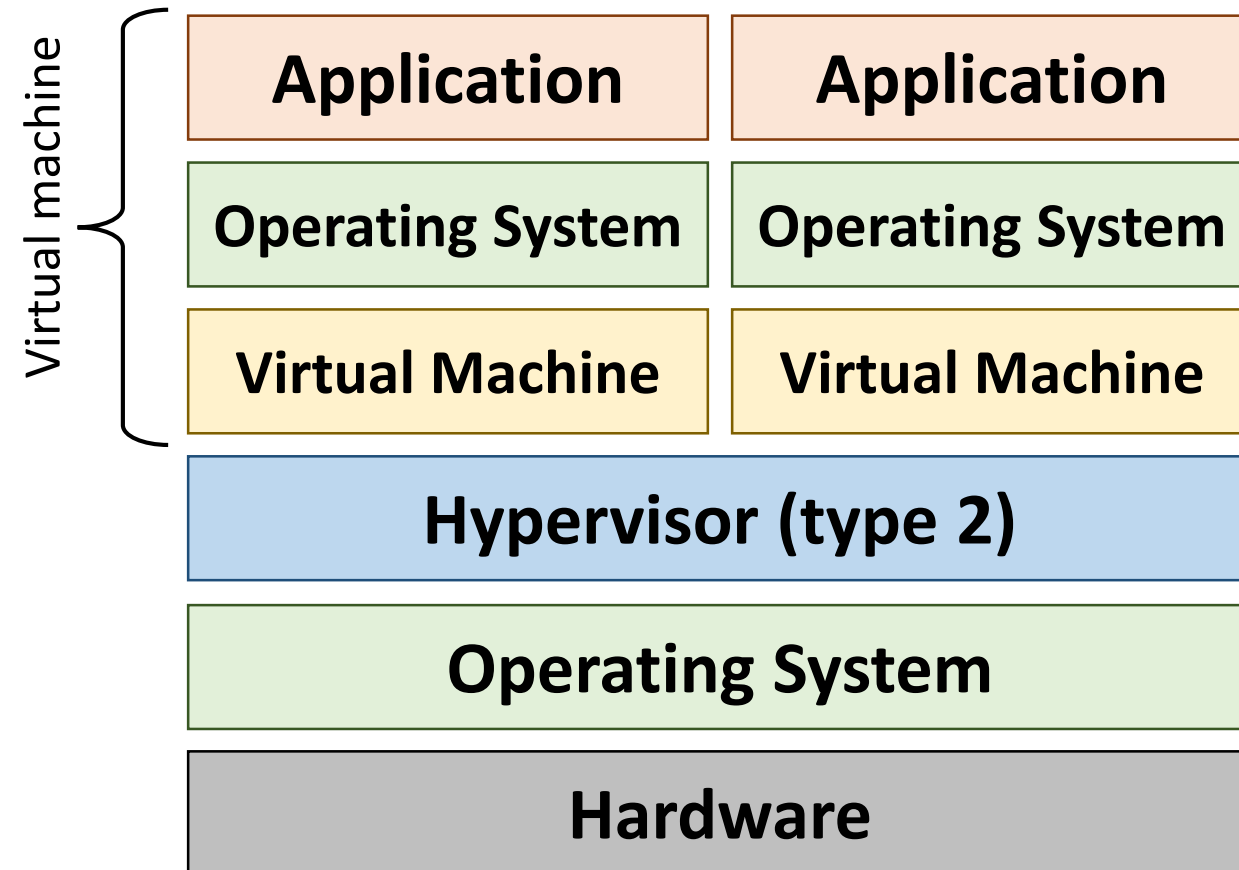
Containers

- Another form of virtualization
- Containers virtualizes Linux
 - Uses LXC (Linux Containers)
- Allows multiple Linux applications to run on a single Linux operating system
 - Reduces or eliminates library conflicts



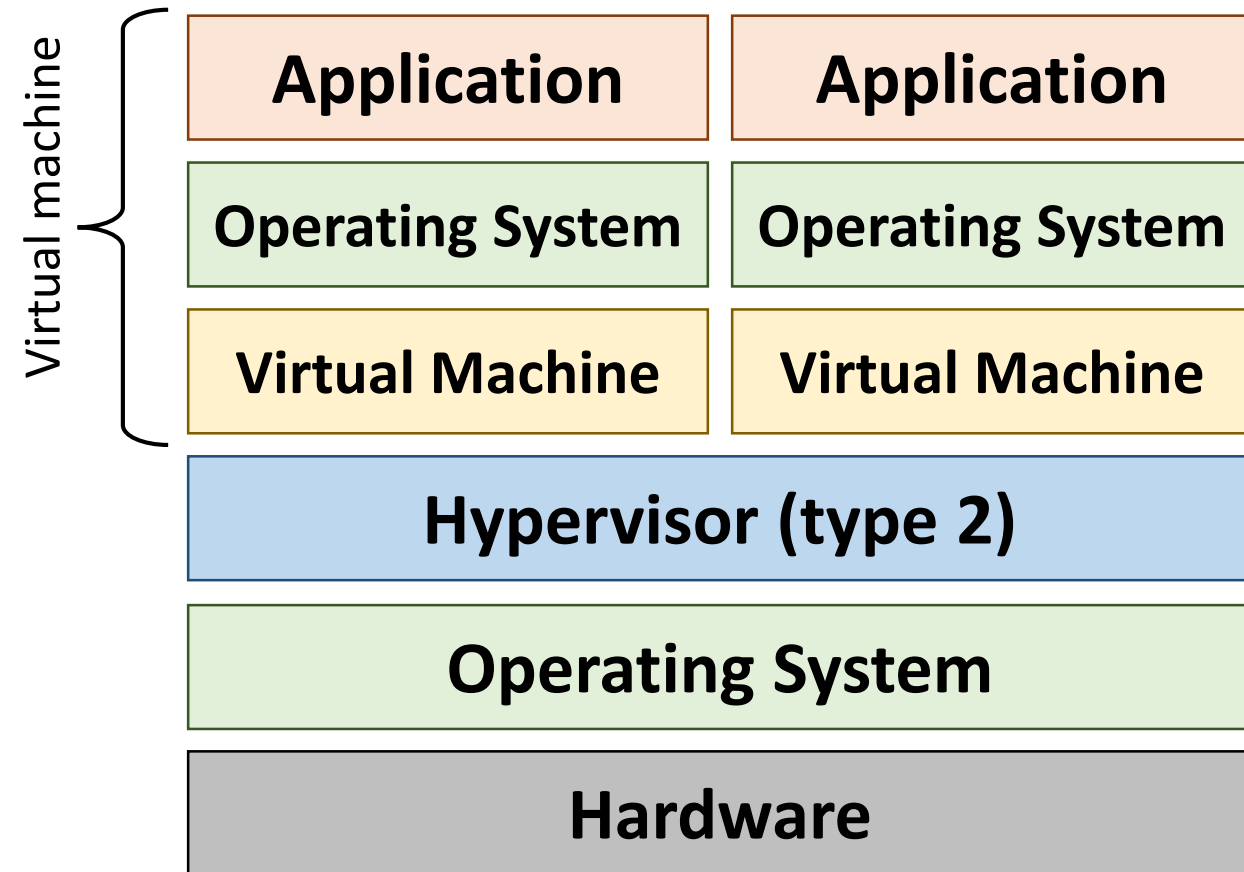
Virtual Machines

VM virtualizes the hardware






Virtual Machines



VM virtualizes the hardware
Container virtualizes Linux





LXC Enabling Technologies

- Namespace
 - Provides isolation
 - Virtualizes Linux resources eg. processes, filesystem, IPC, etc.
- CGroups
 - Allocate resources to the containers
 - Can configure soft and hard limits
- Overlay filesystem
 - Provides a single view of multiple directory by stacking them
 - Provides “copy-on-write” capabilities



Example - Manually Creating Containers

```
ip link add vnet0 type bridge
ip addr add 192.168.0.1/24 dev vnet0
ip link set dev vnet0 up


ip netns add mycntr
ip link add mycntr-veth dev veth peer name mycntr-veth-br

ip link set dev mycntr-veth netns mycntr
ip -n mycntr addr add 192.168.0.10/24 dev mycntr-veth
ip -n mycntr link set dev mycntr-veth up

ip link set dev mycntr-veth-br master vnet0
ip link set dev mycntr-veth-br up

ip -n mycntr route add default via 192.168.0.1


ip netns exec mycntr node main.js
```





Containers and Micro Services

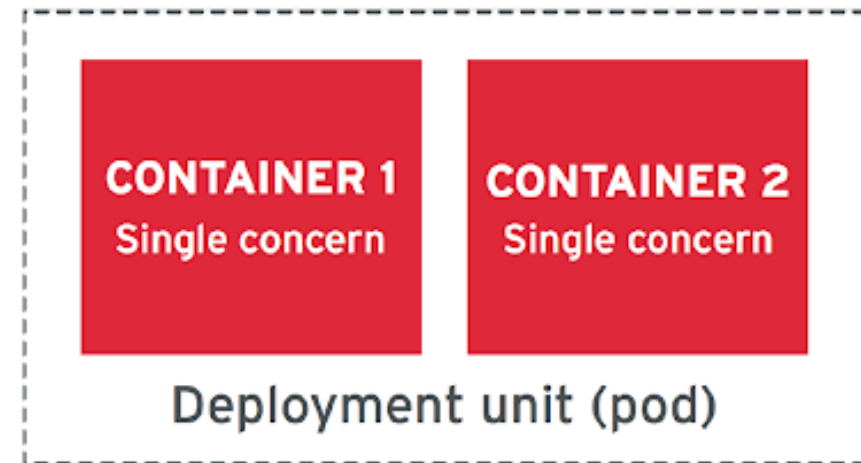
- Containers are extremely lightweight
- Run each service in its own container
- Scale a service by provisioning more of that service
- Communicate with each other via HTTP or queues






Single Concern Principle

- Containers only deliver one service
 - When a micro service is scope to the appropriate granularity
- Treat containers as service primitives
 - Containers interact with each other to deliver higher level service
- Allow for a container be to swapped out in favour of a better implementation of that service
 - Without disrupting the overall service





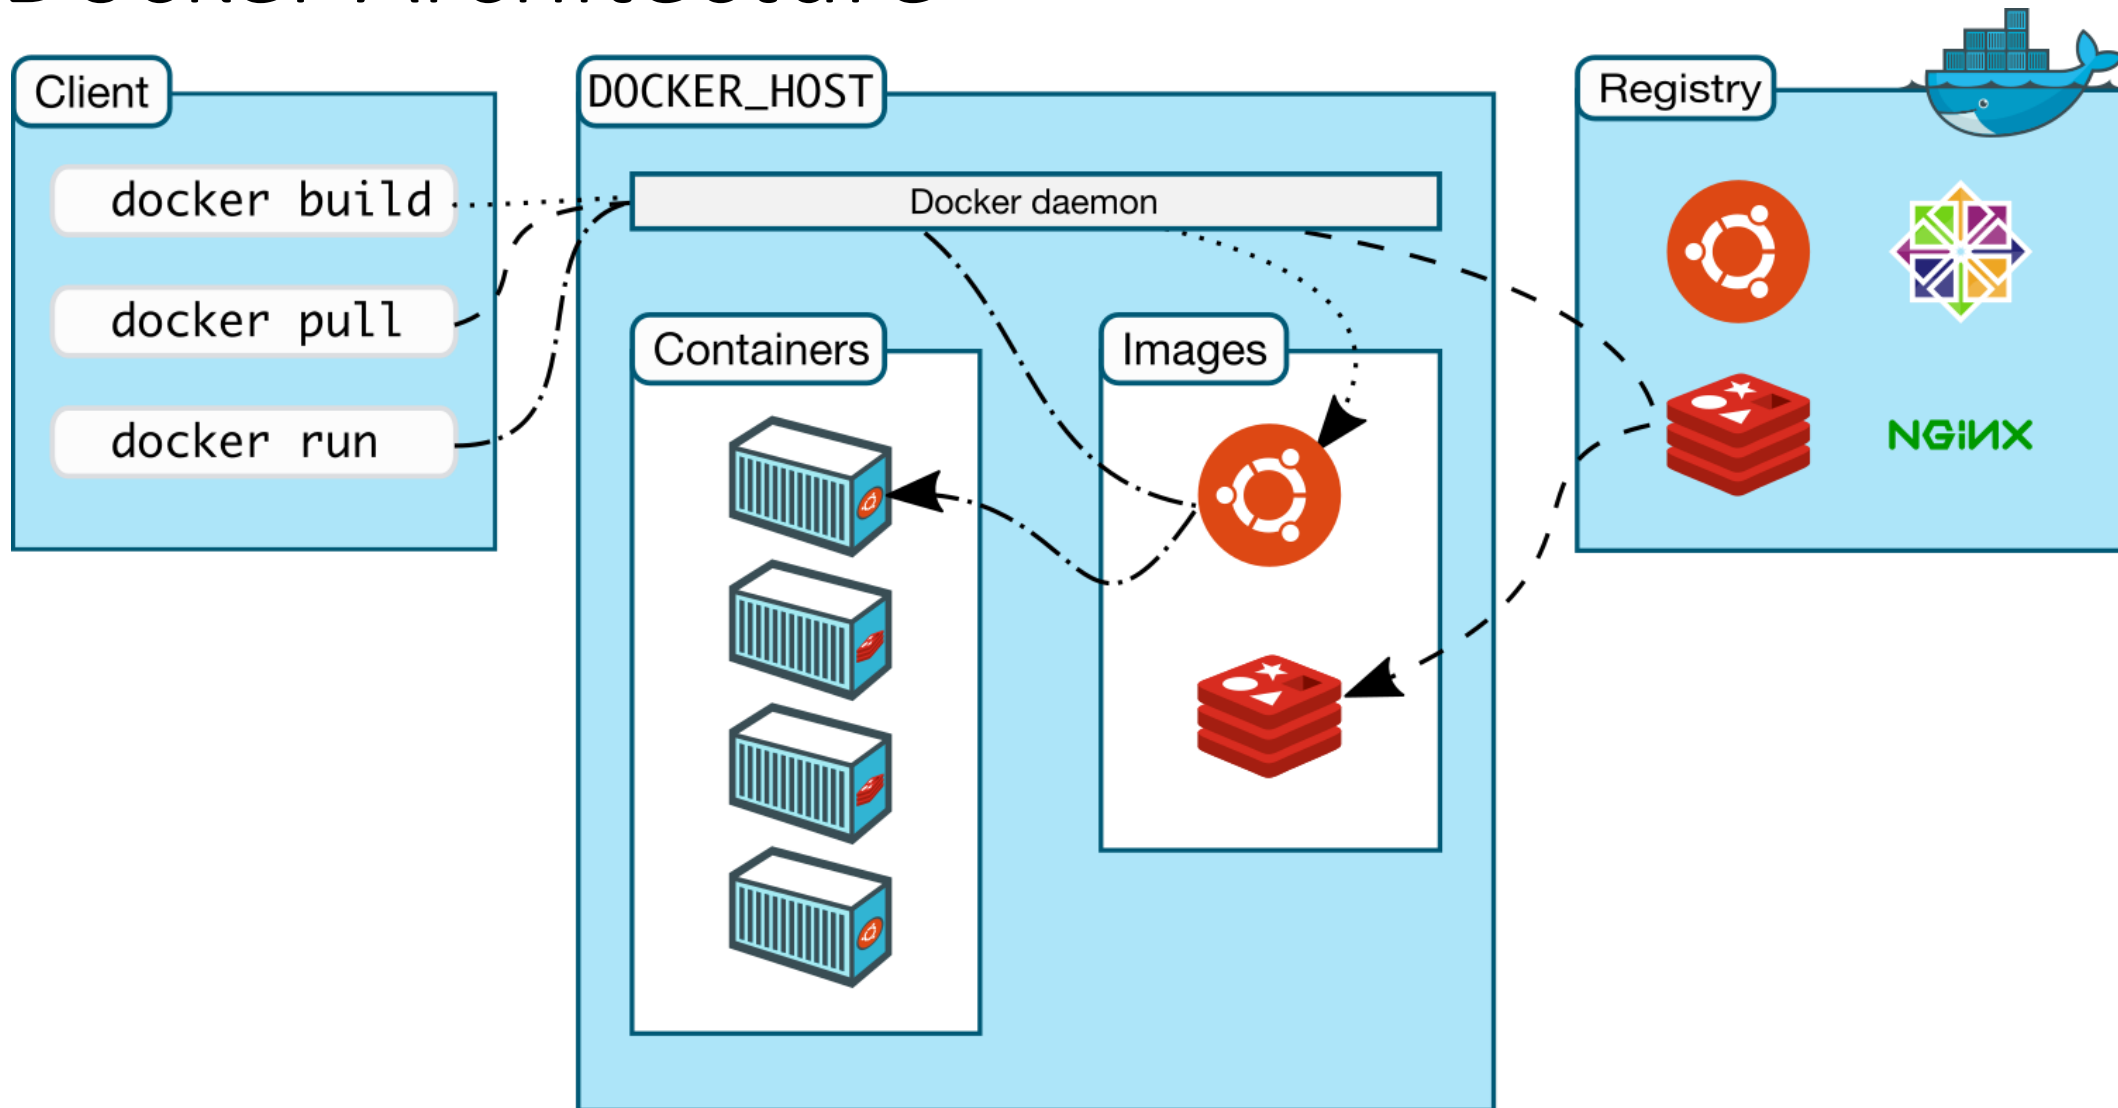
Docker



- A set of LXC management tools for creating “containerized” applications
 - Application isolation by namespaces
 - Specific view of the file system
 - Constrained to a set of resources
- A “image” format

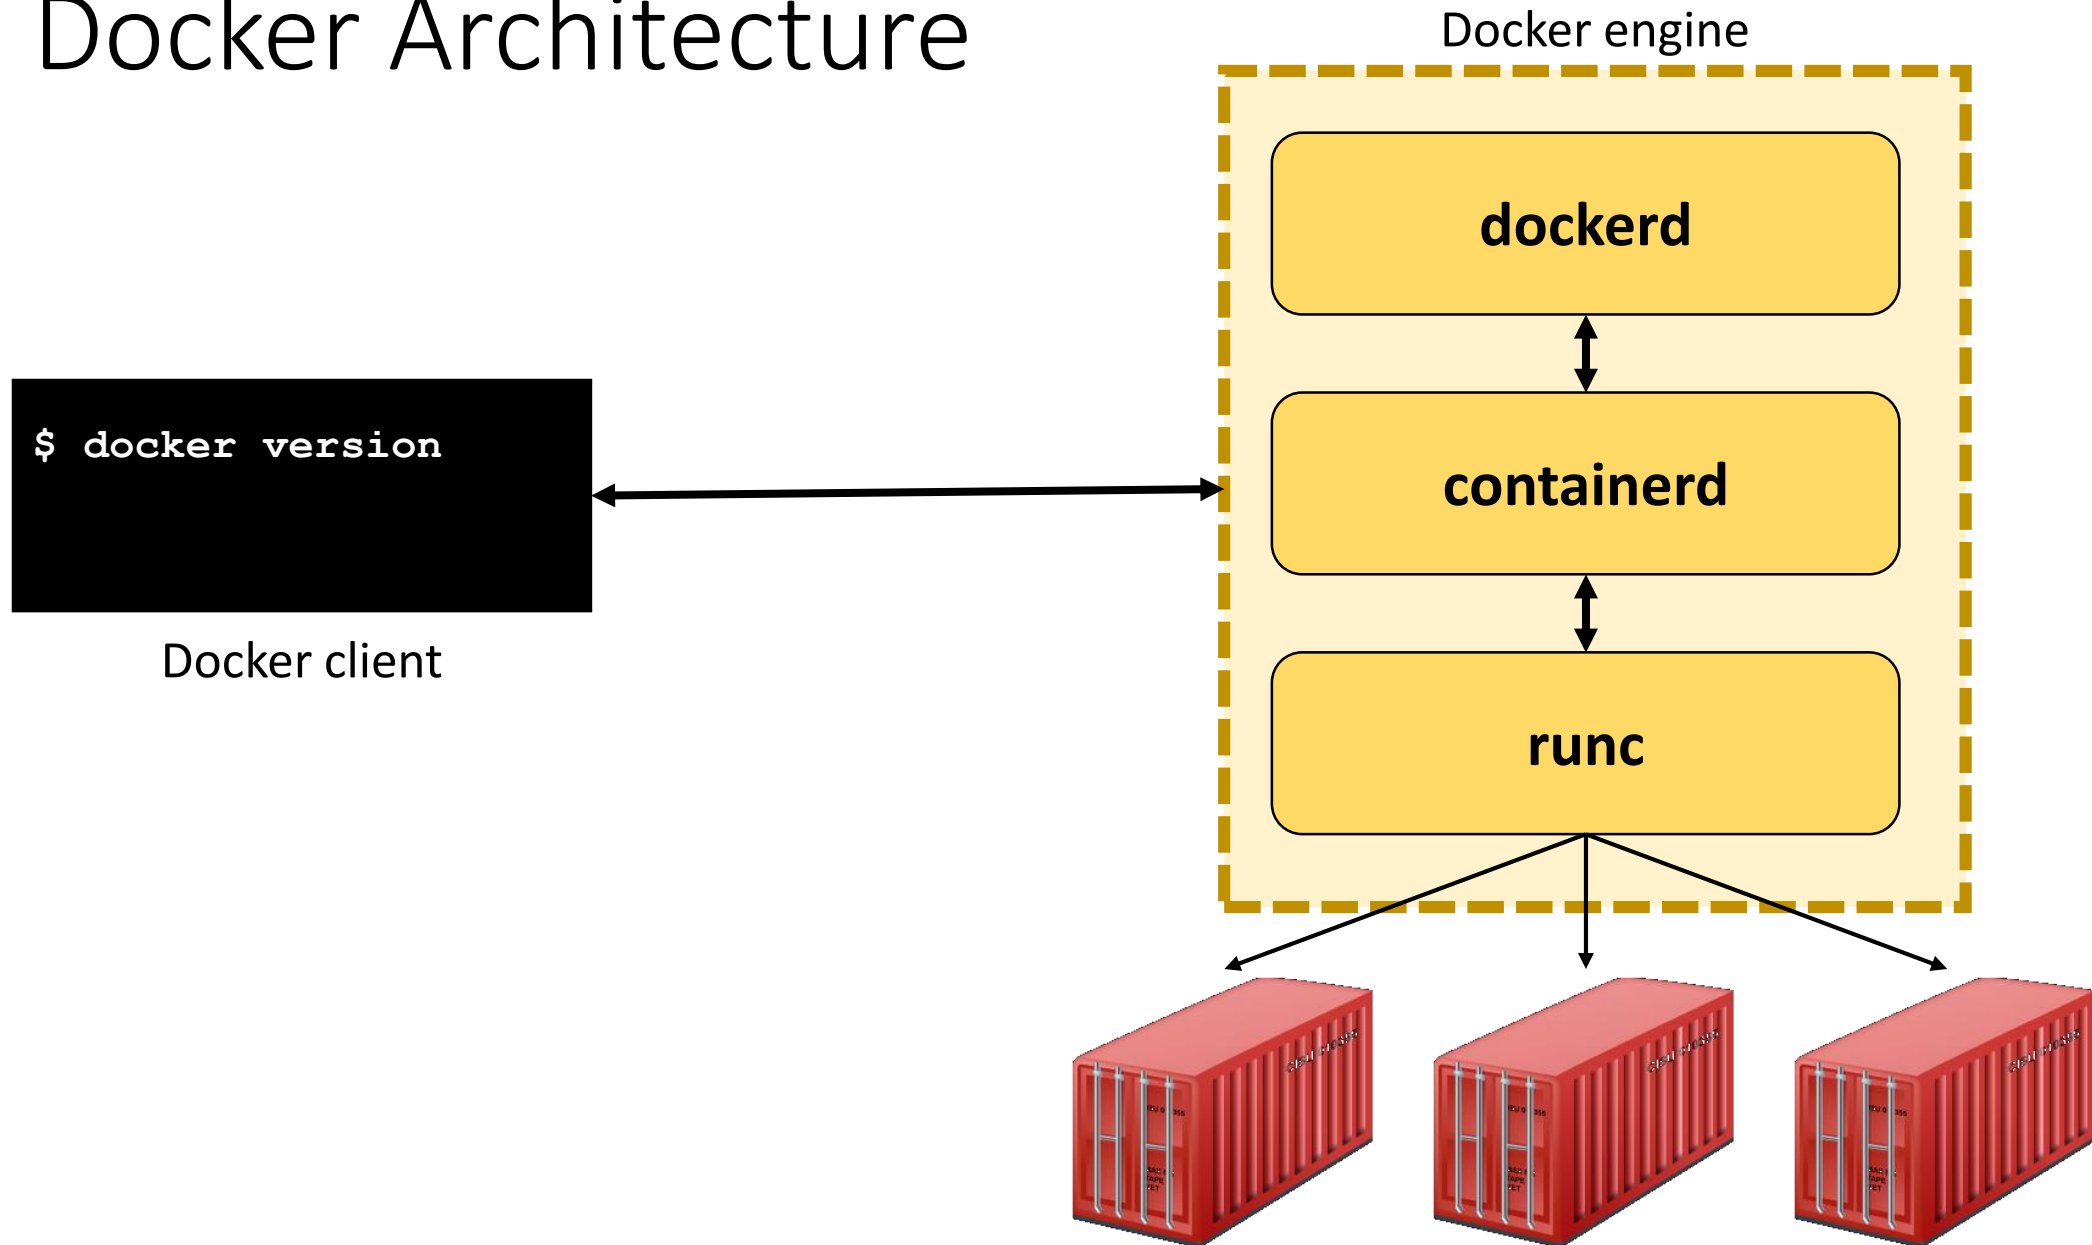


Docker Architecture





Docker Architecture



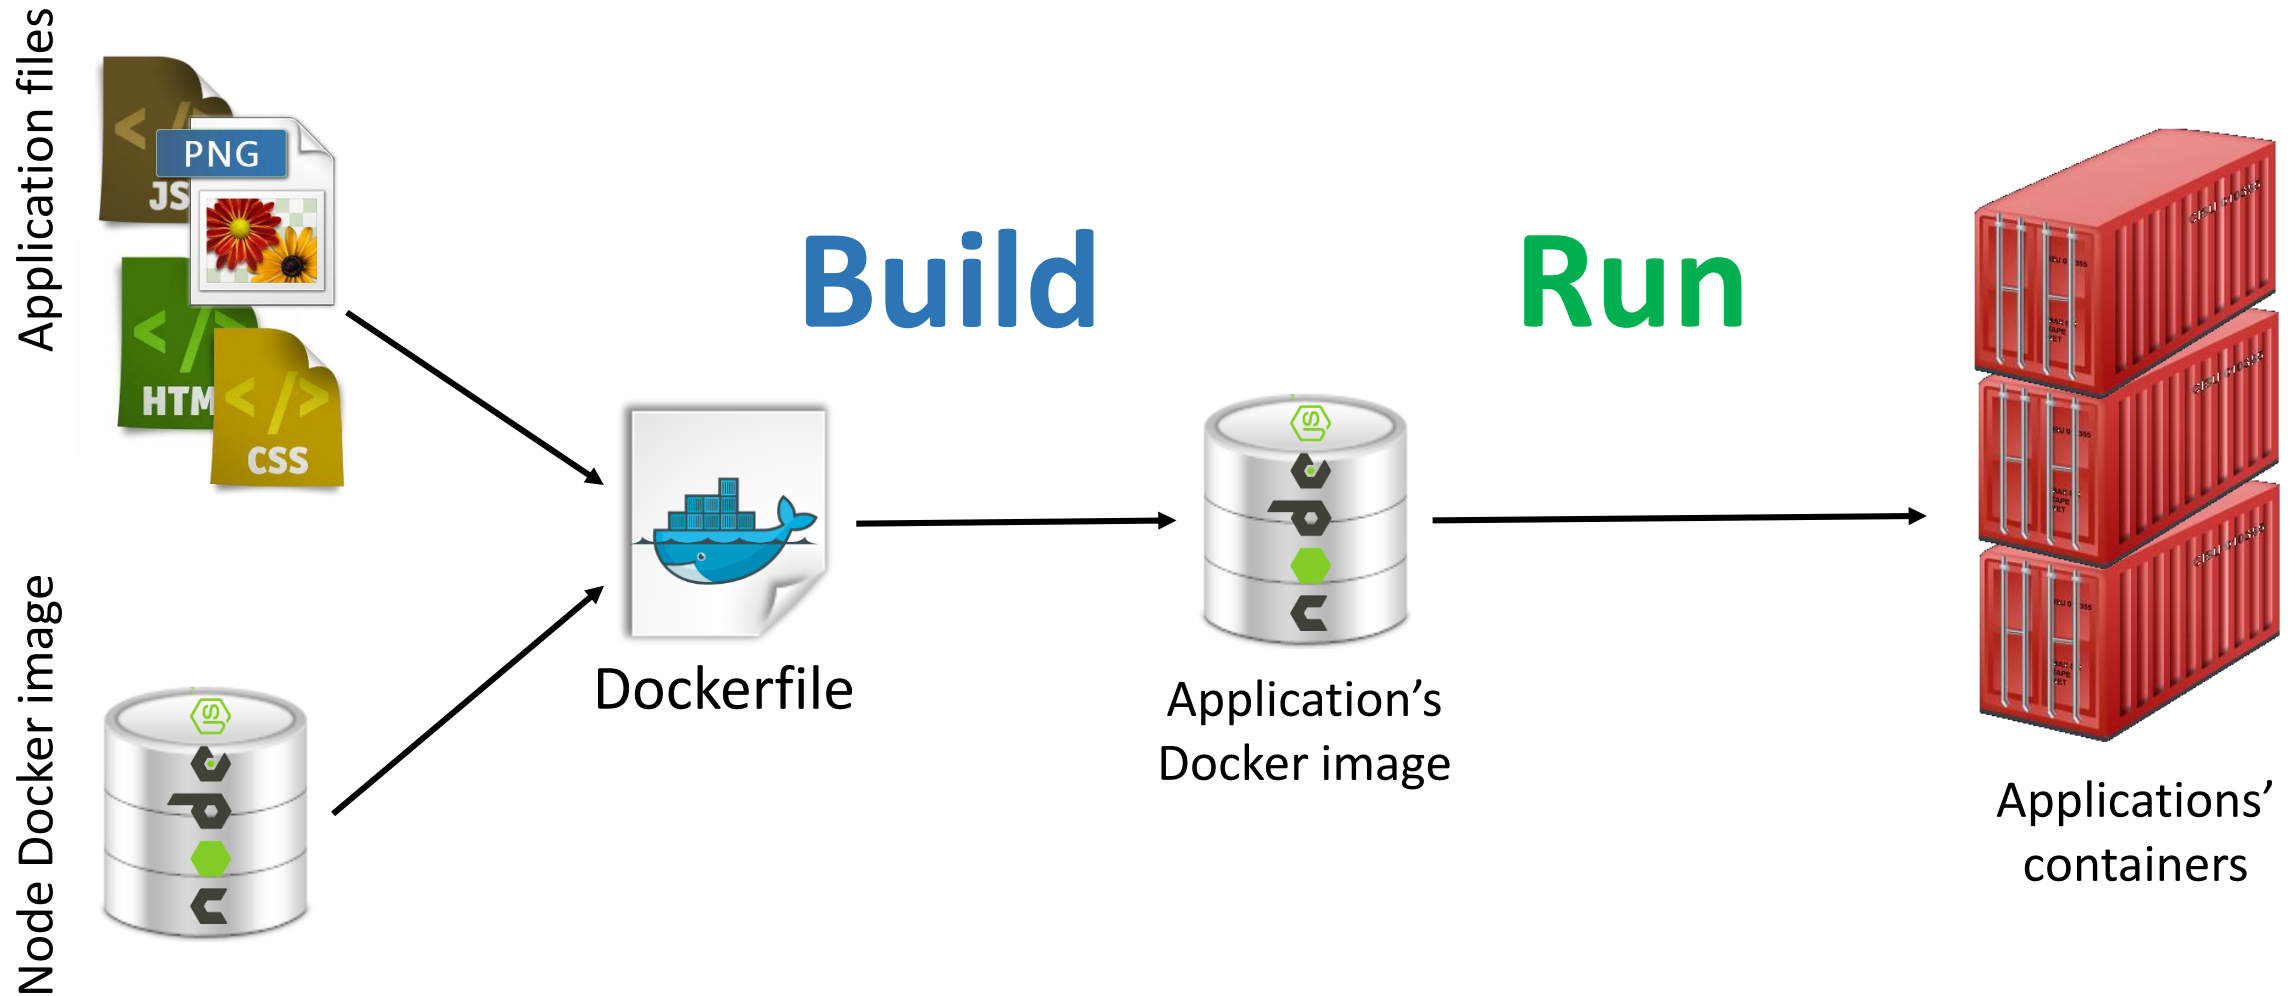


Using Docker

- Tool for packaging, deploying and running applications
- An application and all its dependencies are packaged inside a Docker image
 - Ready to run, like a statically compiled binary
- Application/Docker images are consistent and immutable
 - Will run on everywhere that supports Docker



Docker Workflow





Containerizing an Application

- `Dockerfile` describes how to package an application as a Docker image
 - Like a build file eg. `Makefile`, `pom.xml`
- Describes
 - Application runtime to use
 - Additional packages to install
 - Building the application
 - Executing the application
 - Resources that are needed



Building and Running a Node Application

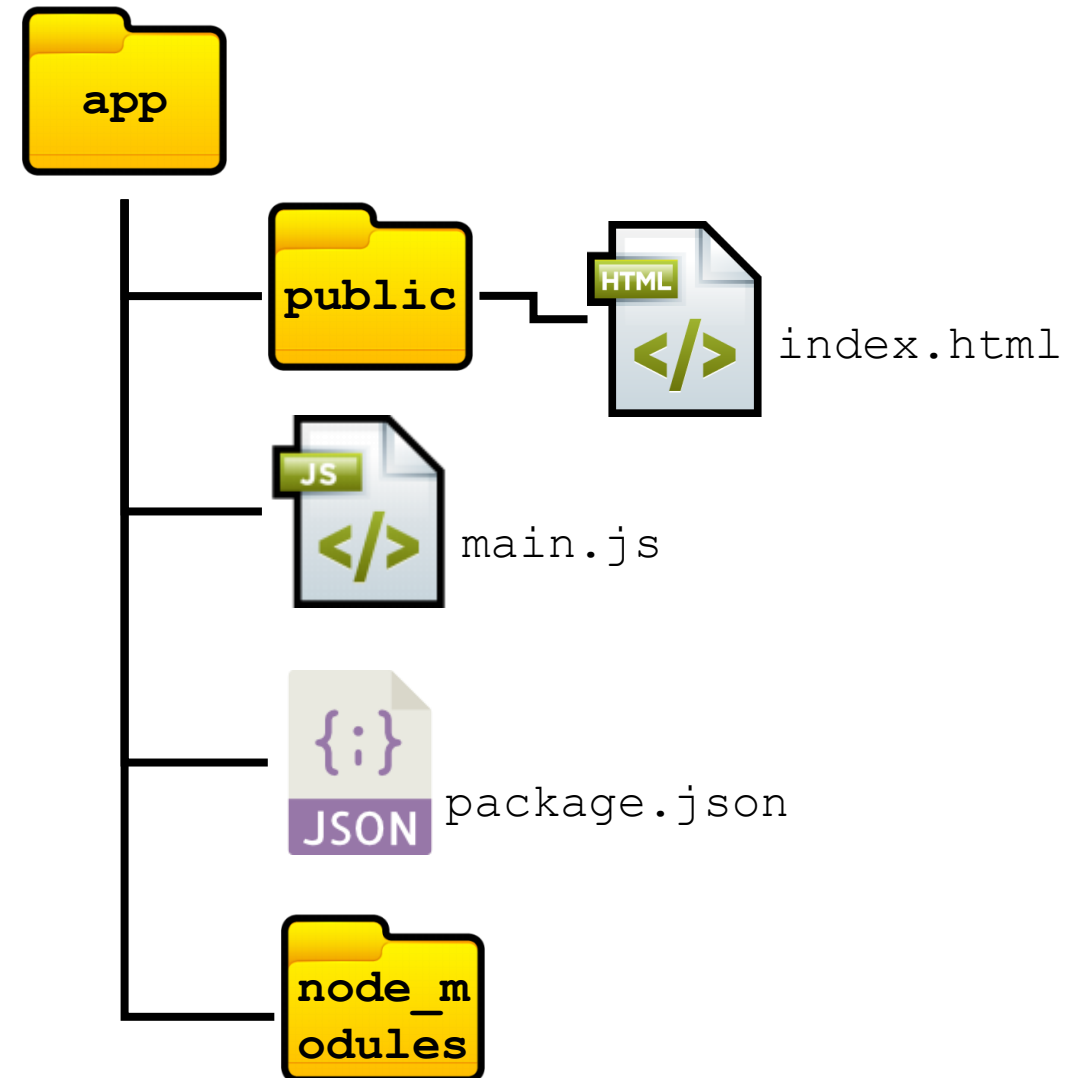
```
mkdir app  
cd app  
mkdir public
```

```
npm init
```

```
npm install --save express
```

```
//edit file main.js  
//edit index.html in public
```

```
node main.js 3000
```





Dockerfile

FROM node:16

Use the node image as the base to build the application

LABEL "name"="myapp"

Add labels to the image

ARG APP_DIR=/app

Build arguments

RUN mkdir \$APP_DIR

Run the command to create a directory

WORKDIR \$APP_DIR

Sets the working directory. Like 'cd' into the directory

Add all these files and directories into \$APP_DIR
Why not node_modules?

ADD main.js .
ADD package.json .
ADD public public

RUN npm ci

Installs dependencies

Tell Docker that the application is listening on \$APP_PORT

ENV APP_PORT=3000

Sets the environment variable

EXPOSE \$APP_PORT

Command to execute when container starts

Provide a default for ENTRYPOINT

ENTRYPOINT ["node", "main.js"]

CMD ["\$APP_PORT"]



Dockerfile

```
FROM node:16

LABEL "name"="myapp"

ARG APP_DIR=/app

RUN mkdir $APP_DIR

WORKDIR $APP_DIR

ADD main.js .
ADD package.json .
ADD public public

RUN npm ci
```

For building
the image

```
ENV APP_PORT=3000
EXPOSE $APP_PORT

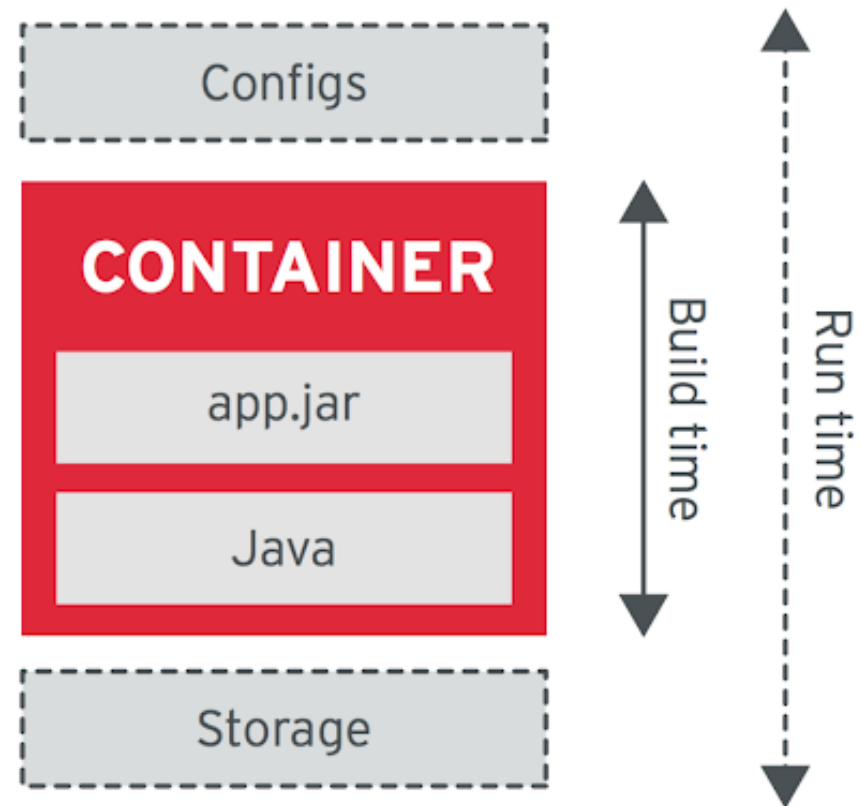
ENTRYPOINT [ "node", "main.js" ]
CMD [ "$APP_PORT" ]
```

For running
the image



Self Containment Principle

- A container should have all dependencies it needs to run the application
 - No other external dependencies
 - Except running on Linux
- Parametrize the things that vary from deployment to deployment
 - Eg. configurations, storage





Building an Image

Application files



Node Docker image




```
docker build --build-arg APP_DIR=/tmp \
-t myapp:v1 .
```

Override the APP_DIR
during build

Image's name

Version

Build context, the
location of the files



Dockerfile

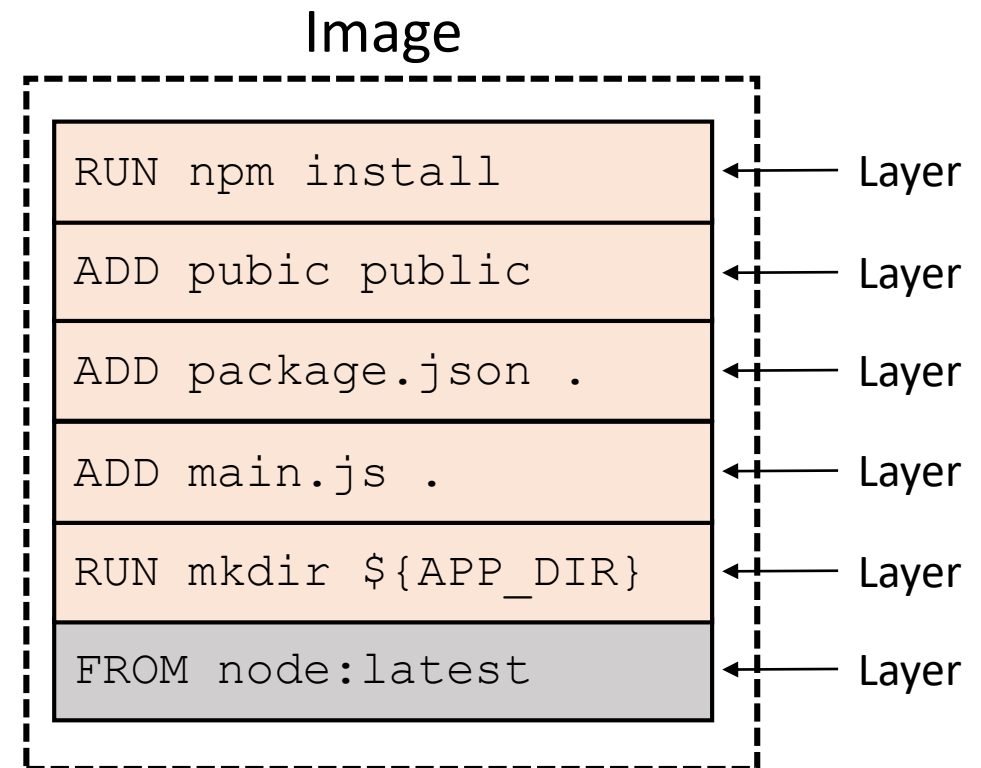


Application's
Docker image



Docker Image

```
docker history myapp:v1
```






Running an Image

Run as daemon
↓
`docker run -d -p 8080:3000 --name app myapp:v1`

Port binding

Container name

Image name



Docker Host: 192.168.0.10

Network traffic to
192.168.0.10:8080 will be routed
to port 3000 in the container

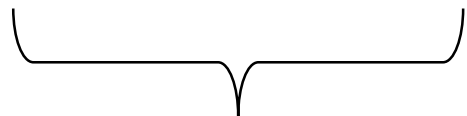




Port Binding

- Container ports are not accessible to the outside world
 - Web applications will not be accessible
 - Will only be accessible to other containers in that network
- Need to specify a port from the host to the container's port
 - Any traffic to the host port will be forwarded to the mapped container port
- Port binding defines this relationship
 - When creating a container

```
docker run -d -p 8080:3000 --name app myapp:v1
```



Port binding



Environment Variables

Set environment variables. Use additional `-e` to set multiple variables

```
docker run -d -p 8080:5000 -e APP_PORT=5000 \
  --name app myapp:v1
```



Container and Image Management

- List all running containers

```
docker ps
```

- Stop a container

```
docker stop mycontainer
```

- Start a container

```
docker start mycontainer
```

- Delete a container

```
docker rm mycontainer
```

- List all images

```
docker image ls
```

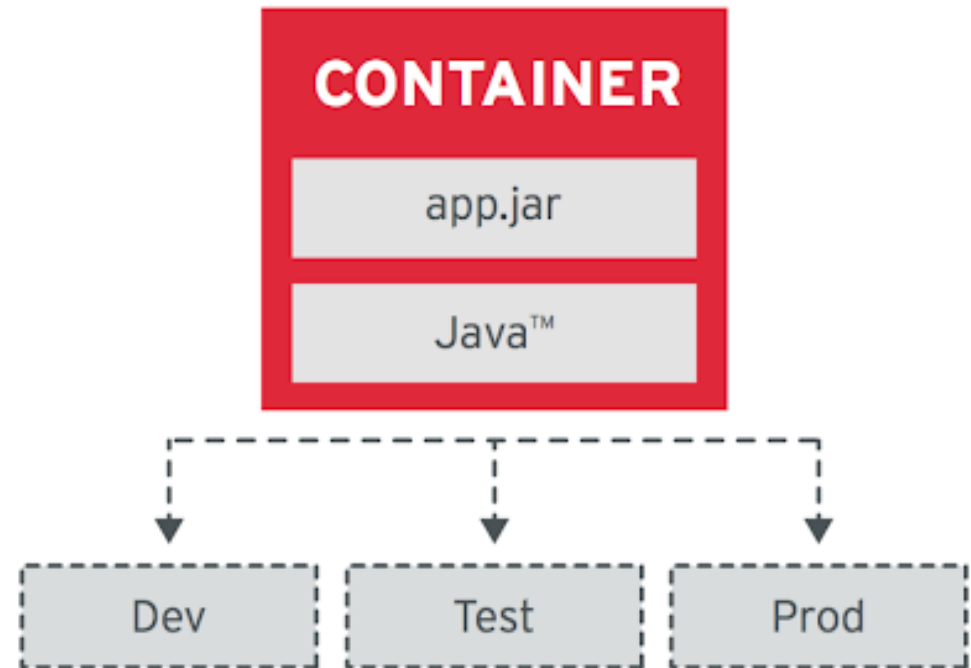
- Delete an image

```
docker rmi myimage
```



Image Immutability Principle

- Images are immutable
- Same image should be use for dev, test and production
 - Only configuration should change
- Do not create snowflakes
 - Exec into a container and patch it
- Should rebuild the image and redeploy





High Observability Principle

- Container are black boxes
- Need to define a standard interface for the container runtime to observe its health
- Suggested observables
 - Readiness - when an application to serve; may be different from when the container is ready
 - Called once at startup
 - Liveness - is the application still alive
 - Called multiple times over the lifetime of the container

- Tracing - allow a request to be traced - OpenTracing
 - Eg. the amount of time spend in a particular database query
- Logs





Lifecycle Conformance Principle

- Receive events from the runtime
 - Inform the container of what is happening
- Application within the container should handle those events





Example of High Observability - Application

```
const pool = mysql.createPool({ ... })
const app = express();
let ready = false;
```

```
app.get('/ready', (req, resp) => {
  resp.status(ready? 200: 400).end();
})
```

Readiness probe. Returns 200 - 399 if the app is ready.
Can double as liveness probe

```
pool.getConnection((err, conn) => {
  conn.ping((err) => {
    ready = !err;
  })
})
```

```
process.on('SIGTERM', () => {
  //Received SIGTERM - perform clean up
})
```

Clean up before the container is removed



Example High Observability - Docker

```
FROM node@sha256:af23..  
...  
HEALTHCHECK --interval=30s --timeout=5s --retries=3 \\  
  CMD curl -s -f http://localhost:${APP_PORT}/ready > /dev/null || exit 1  
ENTRYPOINT [ "node", "main.js" ]  
CMD [ "$APP_PORT" ]
```

Time between health check probe

Number of failed attempts for a container to be considered unhealthy

Returns 0 if successful.
Pass health check

Failed health check



Persistent Data

- Containers are ephemeral
 - Nothing in a container is persisted when a container is removed or dies
 - Eg. Access logs captured by Morgan will not be retained
 - Eg. MySQL database
- Persistent data has to be externalized
 - Written to storage volumes outside of the container
 - When the container is deleted, the data is not delete as well
- Two ways of mapping external storage into Docker
 - Mount a directory from the Docker host into the container
 - Define a Docker volume and mount the volume into the container



Mounting a Local Directory

```
ENV APP_PORT=3000 APP_DIR=/app
```

```
...
```

```
VOLUME ${APP_DIR}/public
```


```
EXPOSE ${APP_PORT}
```

```
...
```

Define a mount point
in the container

```
docker run -d -p 8080:3000 \  
-v /opt/shared:/app/public \  
--name app myapp:v1
```


/app/public /app/public /app/public





Volumes

- Volumes is an abstraction of storage in Docker
 - Different plugins provides storage features
- Properties of volume
 - Local or remote (network attached)
 - Storage type can be block, file or object
 - Block – AWS EBS
 - File – NFS, SMB
 - Object – AWS S3, GCP Cloud Storage





Volume Management

- Create a volume

```
docker volume create myvol
```

- List available volumes

```
docker volume ls
```

- Display the properties of a volume

```
docker volume inspect myvol
```

- Delete a volume

```
docker volume rm myvol
```



Creating and Mounting a Volume

```
docker volume create shared
```

```
docker run -d -p 3000-3100:3000 \  
-v shared:/app/public --name app0 myapp:v1
```

Volume name
without the leading /



Process Disposability Principle


- Containers are ephemeral
 - Can die due to underlying hardware
 - Gets reschedule somewhere else - orchestration
- Externalize your data otherwise its gone
- Design containers to be nimble
 - Quick startup
 - Fail fast





Networking


- Each container get their own
 - Network stack
 - Network interface
 - Virtual network interface (veth)
- Containers connect to their own isolated network
 - Software implementation of 802.1d bridge
- Network are configurable





Docker Network - Bridge


- Allows containers to connect to the same bridge network to communicate
 - Docker creates a default bridge network called `bridge` that all containers are plumbed to if you did not specify any network
 - On Windows bridge is called `nat`





Docker Network - Host

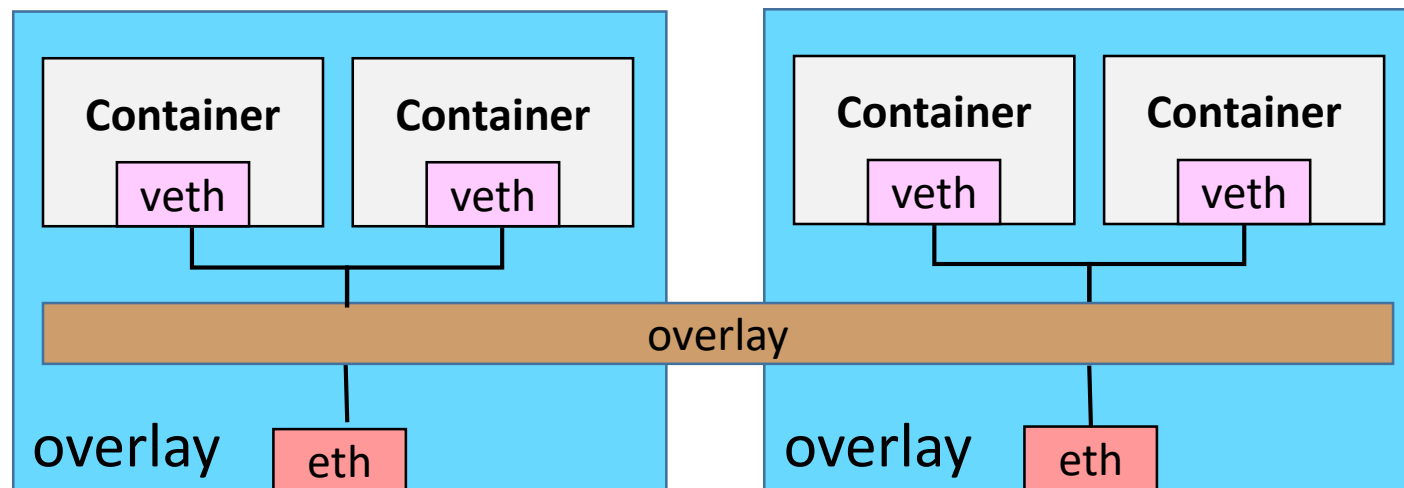
- Container connects into the host's network





Docker Networking - Overlay

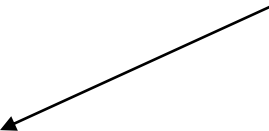
- Allows multiple Docker daemon/host to communicate with each other by creating a network on top of (overlay) of the host network





Attaching to Network

Plumb the container to
bridge network



```
docker run -d -p 8080:3000 --name app myapp:v1
```

```
docker network create -d bridge mynet
```

```
docker run -d -p 8080:3000 --network mynet \  
  --name app myapp:v1
```

```
docker network inspect mynet --format '{{json .Containers}}'
```



Service Discovery

- Docker creates an internal DNS service for User created bridge network
 - Containers connected to the network can communicate via their container name, the `--name` parameter
- Default bridge network, bridge, does not support name resolution via Docker's internal DNS
 - Only user defined bridge networks are supported





Network Management

- Create a network

```
docker network create -d bridge mynet
```

- List available volumes

```
docker network ls
```

- Display network properties

```
docker volume inspect mynet
```

- Delete a volume

```
docker volume rm mynet
```




Deploying Application Stack with Docker

```
docker create network \
  -d bridge app-net

docker create volume data

docker run -d \
  --network app-net \
  --name db northwind-db:v1

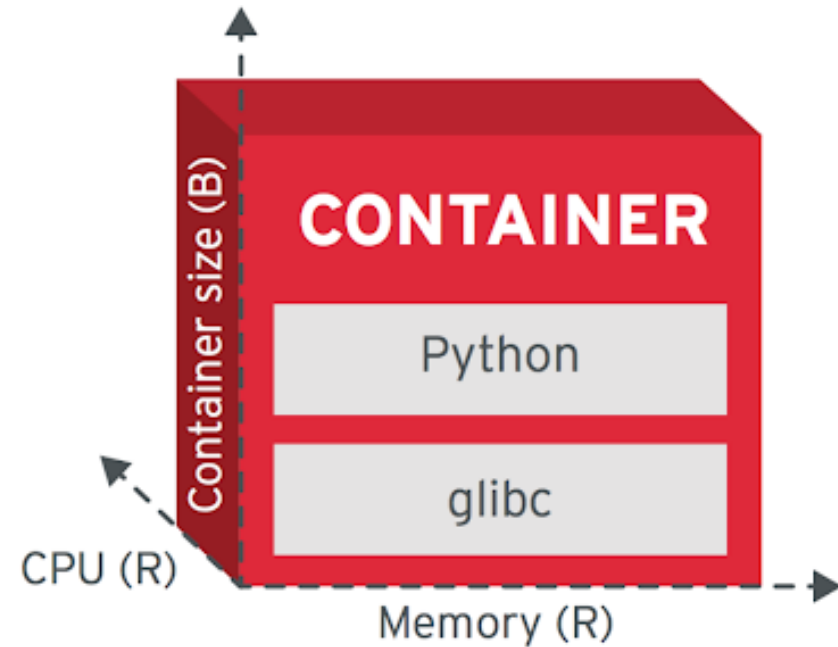
docker run -d -p 8080:3000 \
  -v data:/app/public \
  --network app-net \
  --name app nortwind-app:v1
```





Runtime Confinement Principle

- Many containers may be running on a single host
- Need to sandbox the containers for resource usage
 - Eg. erroneous application don't hog all the resource



```
docker run -d -p 8080:8080 \  
  --cpu-shares=100 \ ← Between 1 - 1024  
  --memory=16m \  
  --blkio-weight=100 \ ← Between 10 - 1000  
  ...
```





Appendix



Docker Networking - Macvlan

- Allows containers to be directly connected to the physical network
 - Each container will have their own IP address
 - Containers appear as independent systems on the physical network







Docker Compose

- Tool for defining and running multi-container application
 - Instead of starting each container individually
- Easily bring up or tear-down entire application stack
- Prioritize resource creation
 - Eg. create networks first before containers
- Docker compose file `docker-compose.yml` consist of the following 3 main parts
 - services – define one or more containers. Each container is considered a service with a name that can be used by other containers for communication
 - networks – define the network to be created
 - volumes – define volumes



docker-compose.yml






docker-compose.yml

```
version: '3'
```

```
volumes:  
  data:
```

```
networks:  
  app-net:
```





docker-compose.yml

...

services:

app:

image: northwind-app:v1

build:

context: ./app

environment:

- APP_PORT=3000
- DB_HOST=db
- DB_USER=root
- DB_PASSWORD=secret

ports:


- 8080-8090:3000

volumes:

- data:/app/public

networks:

- app-net





docker-compose.yml

```
services:
```

```
...
```

```
db:
```

```
  image: northwind-db:v1
```

```
  build:
```


```
    context: ./db
```

```
  environment:
```

```
    - MYSQL_ROOT_PASSWORD=secret
```

```
  networks:
```

```
    - app-net
```





Docker Compose

- Starting a Docker application stack

```
docker-compose up -d
```

- Tearing down a Docker application stack
 - Will remove all containers and network
 - Will not remove volumes and images

```
docker-compose down
```

- Stop the application

```
docker-compose stop
```

- Start the application

```
docker-compose start
```

- Build the images in the stack

```
docker-compose build
```